

Federated Learning Project - Complete Implementation Guide

Project Overview

This project implements federated learning algorithms using different SGD variants, comparing them against the FedAvg baseline.

Team Setup

- **Team size:** 2 people (3 with extra workload)
 - **Deliverables:** 5 federated learning algorithm implementations
-

1. Project Structure

```
federated-learning-project/
├── src/
│   ├── federated_framework.py      # Base framework
│   ├── optimizers.py              # All optimizer implementations
│   ├── models.py                  # CNN and NanoGPT models
│   ├── data_utils.py              # Dataset loading and splitting
│   └── experiments.py            # Experiment runners
└── experiments/
    ├── feminist_experiments.py    # FEMNIST experiments
    └── nanogpt_experiments.py     # NanoGPT experiments
└── results/
    ├── plots/                    # Generated plots
    └── logs/                     # Training logs
└── notebooks/
    └── analysis.ipynb            # Results analysis
└── requirements.txt
└── README.md
```

2. Implementation Strategy

Phase 1: Core Framework (Week 1)

1. Implement base classes (`FederatedOptimizer`, `FederatedClient`, `FederatedServer`)
2. Implement FedAvg baseline
3. Test with simple dataset

Phase 2: Algorithm Implementation (Week 2-3)

Implement 5 algorithm variants (choose from):

Server-Side Optimizers (SGD on clients + optimizer on server):

- FedAdam
- FedAdagrad
- FedRMSprop
- FedMomentum
- FedNAdam
- FedYogi
- FedAdamW

Client-Side Optimizers (Optimizer on clients + averaging on server):

- FedAdam-Client (Adam on clients)
- FedAdagrad-Client
- FedRMSprop-Client
- FedMomentum-Client

Phase 3: Experiments (Week 3-4)

1. Run FEMNIST experiments with CNN
2. Run NanoGPT experiments with custom dataset
3. Test different Dirichlet alpha values (0.1, 0.5, 1.0, 10.0)

Phase 4: Analysis (Week 4)

1. Compare convergence speed
2. Compare final accuracy
3. Analyze communication efficiency
4. Study non-IID robustness

3. Algorithm Selection Guide

Recommended 5 Algorithms:

1. **FedAvg (Baseline) ✓**

- SGD on clients + weighted averaging on server
- Baseline for all comparisons

2. FedAdam-Server

- SGD on clients + Adam on server
- Better convergence in heterogeneous settings

3. FedAdam-Client

- Adam on clients + averaging on server
- Adaptive learning rates help with non-IID data

4. FedMomentum-Server

- SGD on clients + Momentum on server
- Accelerates convergence

5. FedNAdam-Server or FedYogi-Server

- Most sophisticated adaptive method
 - Test state-of-the-art performance
-

4. FEMNIST Dataset Setup

Loading FEMNIST

```
python
from datasets import load_dataset
from torch.utils.data import DataLoader, Subset

# Load FEMNIST
dataset = load_dataset("flwrlabs/femnist", split="train")

# Properties:
# - 62 classes (10 digits + 52 letters, upper and lower case)
# - 28x28 grayscale images
# - Naturally partitioned by writer (user)
```

Non-IID Distribution

The Dirichlet distribution controls data heterogeneity:

- $\alpha = 0.1$: Highly non-IID (each client has few classes)
- $\alpha = 0.5$: Moderately non-IID
- $\alpha = 1.0$: Balanced non-IID
- $\alpha = 10.0$: Nearly IID

```
python
```

```
# Example: Create non-IID split
client_indices = create_non_iid_split(
    dataset,
    num_clients=10,
    alpha=0.5
)
```

5. NanoGPT Setup (Alternative Track)

Creating Custom Dataset

```
python
```

```
# Collect your own text data
texts = [
    "essay1.txt",
    "essay2.txt",
    "journal1.txt",
    # ... different writing samples
]

# Each user's text becomes their local dataset
# This naturally creates a user-private federated setting
```

NanoGPT Model

```
python
```

```

import torch.nn as nn

class NanoGPT(nn.Module):
    def __init__(self, vocab_size, n_embd=384, n_head=6, n_layer=6):
        super().__init__()
        self.token_embedding = nn.Embedding(vocab_size, n_embd)
        self.position_embedding = nn.Embedding(1024, n_embd)
        self.blocks = nn.ModuleList([
            TransformerBlock(n_embd, n_head)
            for _ in range(n_layer)
        ])
        self.ln_f = nn.LayerNorm(n_embd)
        self.lm_head = nn.Linear(n_embd, vocab_size, bias=False)

    def forward(self, idx):
        # Implementation...
        pass

```

6. Experiment Configuration

Standard Hyperparameters

```

python

config = {
    'num_clients': 10,
    'num_rounds': 50,
    'local_epochs': 5,
    'batch_size': 32,
    'learning_rate': 0.01, # Adjust per optimizer
    'alpha': 0.5,         # Dirichlet parameter
    'device': 'cuda'
}

```

Grid Search Suggestions

Test variations:

- **Number of clients:** [5, 10, 20]
- **Local epochs:** [1, 5, 10]
- **Alpha values:** [0.1, 0.5, 1.0]
- **Learning rates:** [0.001, 0.01, 0.1]

7. Evaluation Metrics

Primary Metrics

1. **Test Accuracy:** Final and best accuracy
2. **Test Loss:** Convergence behavior
3. **Convergence Speed:** Rounds to reach target accuracy
4. **Communication Cost:** Total parameters transmitted

Analysis to Include

1. **Learning Curves:** Plot accuracy/loss vs rounds
 2. **Comparative Tables:** Final performance summary
 3. **Ablation Studies:** Effect of α , local epochs, etc.
 4. **Statistical Significance:** T-tests between methods
-

8. Expected Results

Typical Observations

1. **FedAvg:** Good baseline, stable but may be slow
 2. **Adaptive methods (Adam, NAdam):**
 - Faster initial convergence
 - Better with non-IID data
 - May need careful learning rate tuning
 3. **Momentum:**
 - Accelerates convergence
 - Reduces oscillations
 4. **Client-side vs Server-side:**
 - Client-side: More flexible, adapts locally
 - Server-side: Better coordination, less client compute
-

9. Report Structure

Suggested Sections

1. Introduction

- Federated learning overview
- Motivation for SGD variants

2. Background

- FedAvg baseline
- SGD variants description
- Related work

3. Methodology

- Algorithm implementations
- Experimental setup
- Datasets and models

4. Results

- Convergence analysis
- Accuracy comparisons
- Non-IID robustness
- Ablation studies

5. Discussion

- Why certain methods work better
- Trade-offs (speed vs stability)
- Practical recommendations

6. Conclusion

- Summary of findings
- Future work

10. Implementation Tips

Best Practices

1. Reproducibility

```
python
```

```
torch.manual_seed(42)  
np.random.seed(42)
```

2. Logging

```
python
```

```
import wandb # or tensorboard  
wandb.log({  
    'round': round_num,  
    'test_acc': accuracy,  
    'test_loss': loss  
})
```

3. Checkpointing

```
python
```

```
torch.save({  
    'round': round_num,  
    'model_state_dict': model.state_dict(),  
    'optimizer_state': optimizer.state,  
}, f'checkpoint_round_{round_num}.pt')
```

4. Parallel Experiments

```
python
```

```
# Use multiprocessing for multiple configs  
from multiprocessing import Pool  
with Pool(4) as p:  
    results = p.map(run_experiment, configs)
```

Common Pitfalls

✗ **Don't:** Use same learning rate for all optimizers ✓ **Do:** Tune learning rates per optimizer (Adam needs smaller LR)

✗ **Don't:** Ignore client sampling ✓ **Do:** Consider partial client participation

✗ **Don't:** Test only IID data ✓ **Do:** Vary Dirichlet α to test robustness

11. Additional Resources

Papers to Read

1. **FedAvg**: "Communication-Efficient Learning of Deep Networks from Decentralized Data"
2. **FedAdam**: "Adaptive Federated Optimization"
3. **Dirichlet Distribution**: "Measuring the Effects of Non-Identical Data Distribution"

GitHub References

- FedML: <https://github.com/FedML-AI/FedML>
- Flower: <https://github.com/adap/flower>
- PyTorch Examples: <https://github.com/pytorch/examples>

Datasets

- FEMNIST: <https://huggingface.co/datasets/flwrlabs/femnist>
 - LEAF Benchmark: <https://leaf.cmu.edu/>
-

12. Timeline (4 Weeks)

Week 1

- Set up environment
- Implement core framework
- Test FedAvg baseline

Week 2

- Implement 3 additional algorithms
- Verify correctness
- Start FEMNIST experiments

Week 3

- Implement remaining algorithms
- Run comprehensive experiments
- Test different configurations

Week 4

- Analyze results
- Create visualizations
- Write report

13. Deliverables Checklist

- 5 federated learning algorithm implementations
- Working code with documentation
- Experimental results on FEMNIST or NanoGPT
- Comparison plots and tables
- Written report (8-12 pages)
- README with instructions
- (Optional) Presentation slides

Quick Start Commands

```
bash

# Install dependencies
pip install torch torchvision datasets numpy matplotlib pandas tqdm

# Run baseline experiment
python experiments/femnist_experiments.py --algorithm fedavg

# Run comparison
python experiments/femnist_experiments.py --compare-all

# Analyze results
jupyter notebook notebooks/analysis.ipynb
```

Good luck with your project! 